

Assignment 7

CSP 554

Big Data

Ashmita Gupta (A20512498)

Running the demos:

```
gashm@Ashmita MINGW64 ~/OneDrive/Desktop/Big Data
$ ssh -i emr-key-pair-2.pem hadoop@ec2-44-214-182-109.compute-1.amazonaws.com
The authenticity of host 'ec2-44-214-182-109.compute-1.amazonaws.com (44.214.182.109)' can't be established.
ED25519 key fingerprint is SHA256:CSWRLM9knylXsGNWg81HNsDRpoKwKE4iHotHB5HE88w.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-44-214-182-109.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Last login: Thu Oct 12 03:01:33 2023

  _ | _ | _ )
  _ | C | _ /
  _ | \ | _ |
                        Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/

EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRRRR
E::: M::: M::: M::: R:::
EE::: M::: M::: M::: R:::
E:::E EEEEE M::: M::: M::: M::: R::: R:::
E:::E M::: M::: M::: M::: R::: R:::
E:::EEEEEEEE M::: M::: M::: M::: R::: R:::
E::: M::: M::: M::: M::: R::: RR
E:::EEEEEEEE M::: M::: M::: M::: R::: RRRRRR
E:::E M::: M::: M::: M::: R::: R:::
E:::E EEEEE M::: M::: M::: R::: R:::
EE:::EEEEEEEE M::: M::: M::: R::: R:::
E::: M::: M::: M::: R::: R:::
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR RRRRRR

[hadoop@ip-172-31-5-238 ~]$ ls
pydemo.zip sparkdf.zip
[hadoop@ip-172-31-5-238 ~]$ unzip pydemo.zip
Archive:  pydemo.zip
  creating: pydemo/
  inflating: __MACOSX/._pydemo
  inflating: pydemo/testtt.py
  inflating: __MACOSX/pydemo/._testtt.py
  inflating: pydemo/.DS_Store
  inflating: __MACOSX/pydemo/._.DS_Store
  inflating: pydemo/test4.py
  inflating: __MACOSX/pydemo/._test4.py
  inflating: pydemo/twinkle.txt
  inflating: __MACOSX/pydemo/._twinkle.txt
  inflating: pydemo/pydemo.txt
  inflating: __MACOSX/pydemo/._pydemo.txt
  inflating: pydemo/test.py
  inflating: __MACOSX/pydemo/._test.py
  inflating: pydemo/cs595doc2.txt
  inflating: __MACOSX/pydemo/._cs595doc2.txt
  inflating: pydemo/test2.py
  inflating: __MACOSX/pydemo/._test2.py
  inflating: pydemo/test3.py
  inflating: __MACOSX/pydemo/._test3.py
  inflating: pydemo/test3t.py
  inflating: __MACOSX/pydemo/._test3t.py
  inflating: pydemo/twinkle1.py
  inflating: __MACOSX/pydemo/._twinkle1.py
[hadoop@ip-172-31-5-238 ~]$ unzip sparkdf.zip
Archive:  sparkdf.zip
  creating: sparkdf/
  inflating: __MACOSX/._sparkdf
  inflating: sparkdf/dfdemo.txt
  inflating: __MACOSX/sparkdf/._dfdemo.txt
```

```

creating: sparkdf/
inflating: __MACOSX/._sparkdf
inflating: sparkdf/dfdemo.txt
inflating: __MACOSX/sparkdf/._dfdemo.txt
inflating: sparkdf/people.csv
inflating: __MACOSX/sparkdf/._people.csv
inflating: sparkdf/.DS_Store
inflating: __MACOSX/sparkdf/._.DS_Store
inflating: sparkdf/spark3s.py
inflating: __MACOSX/sparkdf/._spark3s.py
inflating: sparkdf/spark2.py
inflating: __MACOSX/sparkdf/._spark2.py
inflating: sparkdf/spark2s.py
inflating: __MACOSX/sparkdf/._spark2s.py
inflating: sparkdf/spark3.py
inflating: __MACOSX/sparkdf/._spark3.py
inflating: sparkdf/spark4s.py
inflating: __MACOSX/sparkdf/._spark4s.py
inflating: sparkdf/spark4.py
inflating: __MACOSX/sparkdf/._spark4.py
inflating: sparkdf/people.txt
inflating: __MACOSX/sparkdf/._people.txt
inflating: sparkdf/spark1.py
inflating: __MACOSX/sparkdf/._spark1.py
inflating: sparkdf/people.json
inflating: __MACOSX/sparkdf/._people.json
inflating: sparkdf/peopleh.csv
inflating: __MACOSX/sparkdf/._peopleh.csv
[hadoop@ip-172-31-5-238 ~]$ ls
__MACOSX  pydemo  pydemo.zip  sparkdf  sparkdf.zip
[hadoop@ip-172-31-5-238 ~]$ cd /home/hadoop/pydemo
[hadoop@ip-172-31-5-238 pydemo]$ hadoop fs -copyFromLocal /home/hadoop/cs595doc2.txt /user/hadoop/cs595doc2.txt
copyFromLocal: /home/hadoop/cs595doc2.txt: No such file or directory
[hadoop@ip-172-31-5-238 pydemo]$ hadoop fs -copyFromLocal /home/hadoop/cs595doc2.txt /user/hadoop/cs595doc2.txt
copyFromLocal: /home/hadoop/cs595doc2.txt: No such file or directory
[hadoop@ip-172-31-5-238 pydemo]$ ls
cs595doc2.txt  pydemo.txt  test2.py  test3.py  test3t.py  test4.py  test.py  testtt.py  twinkle1.py  twinkle.txt
[hadoop@ip-172-31-5-238 pydemo]$ hadoop fs -copyFromLocal /home/hadoop/pydemo/cs595doc2.txt /user/hadoop/cs595doc2.txt
[hadoop@ip-172-31-5-238 pydemo]$ hadoop fs -copyFromLocal /home/hadoop/pydemo/twinkle.txt /user/hadoop/twinkle.txt
[hadoop@ip-172-31-5-238 pydemo]$ pyspark
Python 3.7.16 (default, Aug 30 2023, 20:37:53)
[GCC 7.3.1 20180712 (Red Hat 7.3.1-15)] on linux
Type "help", "copyright", "credits" or "license()" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/10/12 03:15:16 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.
23/10/12 03:15:52 WARN YarnSchedulerBackend$YarnSchedulerEndpoint: Attempted to request executors before the AM has registered!
Welcome to

  ____      _
 / ___|  __| | | |
| |  | |__| | | |
| |  |  __| | | |
| |  | |__| | | |
|_|  |____|_|_|_|

 version 3.4.0-amzn-0

Using Python version 3.7.16 (default, Aug 30 2023 20:37:53)
Spark context Web UI available at http://ip-172-31-5-238.ec2.internal:4040
Spark context available as 'sc' (master = yarn, app id = application_1697079396657_0001).
SparkSession available as 'spark'.
>>> exec(open("/home/hadoop/pydemo/test.py").read())
lines.take(10):
['this is a test of the spark rdd', 'it is a test of pyspark as well', '']
upper.take(10):
['THIS IS A TEST OF THE SPARK RDD', 'IT IS A TEST OF PYSPARK AS WELL', '']
words.take(10):

```

MINGW64:/c/Users/gashm/OneDrive/Desktop/Big Data

FileNotFoundError: [Errno 2] No such file or directory: '/home/hadoop/pydemo/twinkle1.py'

```
>>> exec(open("/home/hadoop/pydemo/twinkle1.py").read())
```

```
['twinkle twinkle little star', 'twinkle twinkle little star']
```

```
>>> exec(open("/home/hadoop/sparkdf/spark1.py").read())
```

```
+-----+
| age | name |
+-----+
| null | Michael |
| 30 | Andy |
| 19 | Justin |
+-----+
```

root

```
-- age: long (nullable = true)
-- name: string (nullable = true)
```

```
>>> exec(open("/home/hadoop/sparkdf/spark2.py").read())
```

```
+-----+
| value |
+-----+
| Michael, 29 |
| Andy, 30 |
| Justin, 19 |
+-----+
```

root

```
-- value: string (nullable = true)
```

```
>>> exec(open("/home/hadoop/sparkdf/spark2s.py").read())
```

Traceback (most recent call last):

File "<stdin>", line 1, in <module>

File "<string>", line 4, in <module>

File "/usr/lib/spark/python/pyspark/sql/readwriter.py", line 602, in text

return self._df(self._jreader.text(self._spark._sc._jvm.PythonUtils.toSeq(paths)))

File "/usr/lib/spark/python/lib/py4j-0.10.9.7-src.zip/py4j/java_gateway.py", line 1323, in __call__

File "/usr/lib/spark/python/pyspark/errors/exceptions/captured.py", line 175, in deco

raise converted from None

pyspark.errors.exceptions.captured.AnalysisException: Column 'age' has a data type of int, which is not supported by Text.

```
>>> exec(open("/home/hadoop/sparkdf/spark3.py").read())
```

```
+-----+
| _c0 | _c1 |
+-----+
| Michael | 29 |
| Andy | 30 |
| Justin | 19 |
+-----+
```

root

```
-- _c0: string (nullable = true)
-- _c1: string (nullable = true)
```

```
>>> exec(open("/home/hadoop/sparkdf/spark3s.py").read())
```

```
+-----+
| name | age |
+-----+
| Michael | 29 |
| Andy | 30 |
| Justin | 19 |
+-----+
```


root

```
-- name: string (nullable = true)
-- age: integer (nullable = true)
```

Exercise 1)

Step A

Start up a EMR cluster as previously, but instead of choosing the “Core Hadoop” configuration chose the “Spark” configuration (see below), otherwise proceed as before.



Create cluster [Info](#)

Name and applications [Info](#)


Name


Amazon EMR release [Info](#)


A release contains a set of applications which can be installed on your cluster.


emr-6.12.0


Application bundle



Spark



Core Hadoop


Flink


HBase


Presto


Trino


Custom

Applications included in bundle

Spark 3.4.0 on Hadoop 3.3.3 YARN with and Zeppelin 0.10.1

AWS Glue Data Catalog settings

Use the AWS Glue Data Catalog to provide an external metastore for your application.

☐ Use for Spark table metadata

Operating system options [Info](#)

☒ Amazon Linux release

☐ Custom Amazon Machine Image (AMI)

☒ Automatically apply latest Amazon Linux updates

Ans: Creating a new cluster with the “Spark” configuration:

[Amazon EMR](#) > [EMR on EC2: Clusters](#) > [Create cluster](#)

Create cluster [Info](#)

Name and applications [Info](#)

Name








My cluster

Amazon EMR release [Info](#)

A release contains a set of applications which can be installed on your cluster.

emr-6.12.0 ▼

Application bundle

Spark 	Core Hadoop 	Flink 	HBase 	Presto 	Trino 	Custom 
---	---	---	---	--	---	--

- ☐ Flink 1.17.0
- ☐ HCatalog 3.1.3
- ☐ Hue 4.11.0
- ☐ Livy 0.7.1
- ☐ Phoenix 5.1.3
- ☒ Spark 3.4.0
- ☐ Tez 0.10.2
- ☐ ZooKeeper 3.5.10

- ☐ Ganglia 3.7.2
- ☐ Hadoop 3.3.3
- ☐ JupyterEnterpriseGateway 2.6.0
- ☐ MXNet 1.9.1
- ☐ Pig 0.17.0
- ☐ Sqoop 1.4.7
- ☐ Trino 414

- ☐ HBase 2.4.17
- ☐ Hive 3.1.3
- ☐ JupyterHub 1.4.1
- ☐ Oozie 5.2.1
- ☐ Presto 0.281
- ☐ TensorFlow 2.11.0
- ☒ Zeppelin 0.10.1

AWS Glue Data Catalog settings

Use the AWS Glue Data Catalog to provide an external metastore for your application.

☐ Use for Spark table metadata

Operating system options [Info](#)

☒ Amazon Linux release

☐ Custom Amazon Machine Image (AMI)

☒ Automatically apply latest Amazon Linux updates

Cluster configuration [Info](#)

Choose a configuration method for the primary, core, and task node groups for your cluster.

☒ Instance groups

Choose one instance type per node group

☐ Instance fleets

Choose any combination of instance types within each node group

Instance groups

Primary

Choose EC2 instance type

m4.xlarge

4 vCore 16 GiB memory EBS only storage

On-Demand price: \$0.200 per instance/hour

Lowest Spot price: \$0.083 (us-east-1e)

Actions ▼

☐ Use multiple primary nodes

To improve cluster availability, use 3 primary nodes with the same configuration and bootstrap actions. You can not use multiple primary nodes with instance fleets.

► Node configuration - *optional*

Core

Remove instance group

Choose EC2 instance type

m4.xlarge

4 vCore 16 GiB memory EBS only storage
On-Demand price: \$0.200 per instance/hour
Lowest Spot price: \$0.083 (us-east-1e)

Actions ▼

► Node configuration - *optional*

Add task instance group

You can add up to 48 more task instance groups.

► EBS root volume - *optional*

Cluster scaling and provisioning option [Info](#)

Set up scaling and provisioning configurations for the core and task node groups for your cluster.

Choose an option

☒ Set cluster size manually

Use this option if you know your workload patterns in advance.

☐ Use EMR-managed scaling

Monitor key workload metrics so that EMR can optimize the cluster size and resource utilization.

☐ Use custom automatic scaling

To programmatically scale core and task nodes, create custom automatic scaling policies.

Provisioning configuration

Set the size of your core instance group. Amazon EMR attempts to provision this capacity when you launch your cluster.

Creating a new key pair for the cluster:

[EC2](#) > [Key pairs](#) > Create key pair

Create key pair [Info](#)

Key pair

A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.

Name

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type [Info](#)

☒ RSA☐ ED25519

Private key file format

☒ .pem
For use with OpenSSH☐ .ppk
For use with PuTTY

Tags - *optional*

No tags associated with the resource.

You can add up to 50 more tags.

[Cancel](#)

[Create key pair](#)

Cluster created:

The screenshot shows the AWS Management Console interface for an Amazon EMR cluster. At the top, a green banner states: "Your cluster 'My cluster' has been successfully created." Below this, the breadcrumb navigation shows "Amazon EMR > EMR on EC2: Clusters > My cluster". The main heading is "My cluster", with a timestamp "Updated less than a minute ago" and buttons for "Terminate", "Clone in AWS CLI", and "Clone".

The "Summary" section is divided into four columns:

- Cluster info:** Cluster ID is j-258PDY1V725YB. Cluster configuration is Instance groups. Capacity is 1 Primary, 1 Core, 0 Task.
- Applications:** Amazon EMR version is emr-6.12.0. Installed applications are Spark 3.4.0, Zeppelin 0.10.1.
- Cluster management:** Log destination in Amazon S3 is aws-logs-020428218454-us-east-1/elasticmapreduce. Persistent application Uls are Spark History Server and YARN timeline server. Primary node public DNS is ec2-44-214-182-109.compute-1.amazonaws.com. A link to "Connect to the Primary Node using SSH" is provided.
- Status and time:** Status is Waiting. Creation time is October 11, 2023, 21:51 (UTC-05:00). Elapsed time is 6 minutes, 6 seconds.

Below the summary, there are tabs for "Properties", "Bootstrap actions", "Instances (Hardware)", "Steps", "Applications", "Configurations", "Monitoring", "Events", and "Tags (1)". The "Properties" tab is active, showing "Operating system" (Amazon Linux release 2.0.20230906.0), "Cluster logs" (Archive log files to Amazon S3, Turned on, Amazon S3 location), and "Cluster termination" (Edit cluster termination, Termination option: Automatically terminate cluster after idle time).

ssh to cluster:

```
gashm@Ashmita MINGW64 ~/OneDrive/Desktop/Big Data
$ ssh -i emr-key-pair-2.pem hadoop@ec2-44-214-182-109.compute-1.amazonaws.com
The authenticity of host 'ec2-44-214-182-109.compute-1.amazonaws.com (44.214.182.109)' can't be established.
ED25519 key fingerprint is SHA256:CSwRLM9kny1XsGNWg81HnsDRpKwke4iH0tHB5HE88w.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-44-214-182-109.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Last login: Thu Oct 12 03:01:33 2023
```

```
  _ | _ | _ )
 _ | ( _ | /
 _ | \ _ | _ |
          Amazon Linux 2 AMI
```

```
https://aws.amazon.com/amazon-linux-2/
```

```
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRR
E:::EEEEEEEEEEEEEEEE M:::M M:::M R:::R
EE:::EEEEEEEEEEEEEEEE M:::M M:::M R:::RRRRRRRRRRRR
E:::E EEEEE M:::M M:::M RR::R R:::R
E:::E M:::M M:::M M:::M R::R R:::R
E:::EEEEEEEEEEEE M:::M M:::M M:::M R::RRRRRRRRRRRR
E:::EEEEEEEEEEEE M:::M M:::M M:::M R::RRRRRRRRRRRR
E:::EEEEEEEEEEEE M:::M M:::M M:::M R::RRRRRRRRRRRR
E:::E M:::M M:::M M:::M R::R R:::R
E:::E EEEEE M:::M MMM M:::M R::R R:::R
EE:::EEEEEEEEEEEE M:::M M:::M R::R R:::R
E:::EEEEEEEEEEEE M:::M M:::M RR::R R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRR RRRRRR
```

```
[hadoop@ip-172-31-5-238 ~]$
```

Step B

Use the TestDataGen program from previous assignments to generate new data files.

Copy both generated files to the HDFS directory “/user/hadoop”

Ans. Copied both generated files to the HDFS directory “/user/hadoop”

```
[hadoop@ip-172-31-5-238 ~]$ ls
__MACOSX  pydemo  pydemo.zip  sparkdf  sparkdf.zip  TestDataGen.class
[hadoop@ip-172-31-5-238 ~]$ java TestDataGen
Magic Number = 205356
[hadoop@ip-172-31-5-238 ~]$ ls
foodplaces205356.txt  foodratings205356.txt  __MACOSX  pydemo  pydemo.zip  sparkdf  sparkdf.zip  TestDataGen.class
[hadoop@ip-172-31-5-238 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodplaces205356.txt /user/hadoop/foodplaces205356.txt
[hadoop@ip-172-31-5-238 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodratings205356.txt /user/hadoop/foodratings205356.txt
```

Magic number: 205356

Step C

Load the ‘foodratings’ file as a ‘csv’ file into a DataFrame called foodratings. When doing so specify a schema having fields of the following names and types:

Field Name	Field Type
Name	String
food1	Integer
food2	Integer
food3	Integer
food4	Integer
placeid	Integer

As the results of this exercise provide the magic number, *the code you execute* and screen shots of the following commands:

```
foodratings.printSchema()
```

```
foodratings.show(5)
```

Ans.

Loading the ‘foodratings’ file as a ‘csv’ file into a DataFrame called foodratings:

```
>>> from pyspark.sql.types import *
>>> struct1 = StructType(
...     [
...         StructField("name", StringType(), True),
...         StructField("food1", IntegerType(), True),
...         StructField("food2", IntegerType(), True),
...         StructField("food3", IntegerType(), True),
...         StructField("food4", IntegerType(), True),
...         StructField("placeid", IntegerType(), True),
...     ]
... )
>>> foodratings = spark.read.schema(struct1).csv('/user/hadoop/foodratings205356.txt')
```

```
foodratings.printSchema()
```

```
foodratings.show(5)
```

```
>>> foodratings.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings.show(5)
+----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+----+-----+-----+-----+-----+
| Joe|    8|   12|    1|   29|      2|
| Mel|   31|   34|   44|   27|      3|
| Joe|   13|   41|   12|   10|      4|
| Jill|  23|    6|    3|   45|      3|
| Joe|   33|   16|   28|   28|      2|
+----+-----+-----+-----+
only showing top 5 rows
```

Exercise 2)

Load the 'foodplaces' file as a 'csv' file into a DataFrame called foodplaces. When doing so specify a schema having fields of the following names and types:

Field Name	Field Type
placeid	Integer
placename	String

As the results of this exercise provide *the code you execute* and screen shots of the following commands:

```
foodplaces.printSchema()
```

```
foodplaces.show(5)
```

Ans.

Load the 'foodplaces' file as a 'csv' file into a DataFrame called foodplaces :

```
>>> struct2 = StructType().add("placeid", IntegerType(), True).add("placename", StringType(), True)
>>> foodplaces = spark.read.schema(struct2).csv('/user/hadoop/foodplaces205356.txt')
>>> foodplaces.printSchema()
```

Magic number: **205356**

```
foodplaces.printSchema()
```

```
foodplaces.show(5)
```

```
>>> foodplaces.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> foodplaces.show(5)
+-----+-----+
|placeid| placename|
+-----+-----+
|      1| China Bistro|
|      2|   Atlantic|
|      3|  Food Town|
|      4|   Jake's|
|      5|  Soup Bowl|
+-----+-----+
```

Exercise 3)

Step A

Register the DataFrames created in exercise 1 and 2 as tables called “foodratingsT” and “foodplacesT”

Ans.

Registering the DataFrames as tables called “foodratingsT” and “foodplacesT”:

```
>>> foodratings.createOrReplaceTempView("foodratingsT")
>>> foodplaces.createOrReplaceTempView("foodplacesT")
>>> foodratings_ex3a = spark.sql("select * from foodratingsT where food2 < 25 and food4 > 40")
```

Step B

Use a SQL query on the table “foodratingsT” to create a new DataFrame called foodratings_ex3a holding records which meet the following condition: food2 < 25 and food4 > 40. Remember, when defining conditions in your code use maximum parentheses.

As the results of this step *provide the code you execute* and screen shots of the following commands:

```
foodratings_ex3a.printSchema()
```

```
foodratings_ex3a.show(5)
```

Ans.

foodratings_ex3a holding records which meet the condition: food2 < 25 and food4 > 40:

```
>>> foodratings_ex3a = spark.sql("select * from foodratingsT where food2 < 25 and food4 > 40")
```

```
foodratings_ex3a.printSchema()
```

```
foodratings_ex3a.show(5)
```

```
>>> foodratings_ex3a.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings_ex3a.show(5)
+----+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+----+-----+-----+-----+-----+-----+
|Jill|  23|   6|   3|  45|     3|
| Sam|  50|   5|  16|  50|     1|
| Joe|  26|  17|  32|  43|     2|
| Joe|  26|   8|  40|  46|     5|
|Jill|  49|  23|  25|  48|     5|
+----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Step C

Use a SQL query on the table "foodplacesT" to create a new DataFrame called foodplaces_ex3b holding records which meet the following condition: placeid > 3

As the results of this step *provide the code you execute* and screen shots of the following commands:

```
foodplaces_ex3b.printSchema()
```

```
foodplaces_ex3b.show(5)
```

Ans.

Below is the SQL query on the table "foodplacesT" to create a new DataFrame called foodplaces_ex3b holding records which meet the following condition: placeid > 3

```
>>> foodplaces_ex3b = spark.sql("select * from foodplacesT where placeid >3")
```

```
foodplaces_ex3b.printSchema()
```

```
foodplaces_ex3b.show(5)
```

```
>>> foodplaces_ex3b.printSchema()
root
|-- placeid: integer (nullable = true)
|-- placename: string (nullable = true)

>>> foodplaces_ex3b.show(5)
+-----+-----+
|placeid|placename|
+-----+-----+
|      4|  Jake's |
|      5| Soup Bowl|
+-----+-----+
```

Exercise 4)

Use a transformation (not a SparkSQL query) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex4 that includes only those records (rows) where the 'name' field is "Mel" and food3 < 25.

As the results of this step provide the code you execute and screen shots of the following commands:

```
foodratings_ex4.printSchema()
```

```
foodratings_ex4.show(5)
```

Ans.

New DataFrame called foodratings_ex4 includes only those records where the 'name' field is "Mel" and food3 < 25:

```
>>> foodratings_ex4 = foodratings.filter((foodratings['name'] == "Mel") & (foodratings['food3'] < 25))
>>> foodratings_ex4.printSchema()
root
|-- name: string (nullable = true)
|-- food1: integer (nullable = true)
|-- food2: integer (nullable = true)
|-- food3: integer (nullable = true)
|-- food4: integer (nullable = true)
|-- placeid: integer (nullable = true)

>>> foodratings_ex4.show(5)
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
| Mel|    6|   49|   11|   47|      3|
| Mel|    3|   19|   16|   40|      2|
| Mel|   31|    1|    9|   13|      1|
| Mel|   17|    4|   11|   22|      5|
| Mel|   41|   44|    1|   40|      1|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Exercise 5)

Use a transformation (**not a SparkSQL query**) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex5 that includes only the columns (fields) 'name' and 'placeid'

As the results of this step provide the code you execute and screen shots of the following commands:

```
foodratings_ex5.printSchema()
```

```
foodratings_ex5.show(5)
```

Ans.

New DataFrame called foodratings_ex5 that includes only the columns (fields) 'name' and 'placeid':

```
>>> foodratings_ex5 = foodratings.select(foodratings['name'], foodratings['placeid'])
>>> foodratings_ex5.printSchema()
root
 |-- name: string (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings_ex5.show(5)
+-----+-----+
|name|placeid|
+-----+-----+
| Joe|      2|
| Mel|      3|
| Joe|      4|
| Jill|     3|
| Joe|      2|
+-----+-----+
only showing top 5 rows
```

Exercise 6)

Use a transformation (**not a SparkSQL query**) to create a new DataFrame called ex6 which is the inner join, on placeid, of the DataFrames 'foodratings' and 'foodplaces' created in exercises 1 and 2

As the results of this step provide the code you execute and screen shots of the following commands:

```
ex6.printSchema()
```

```
ex6.show(5)
```

Ans.

New DataFrame called ex6 which is the inner join, on placeid, of the DataFrames 'foodratings' and 'foodplaces' created in exercises 1 and 2:

```
>>> ex6 = foodratings.join(foodplaces, foodratings.placeid == foodplaces.placeid, 'inner')
>>> ex6.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> ex6.show(5)
+-----+-----+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|placeid|placename|
+-----+-----+-----+-----+-----+-----+-----+
| Joe|    8|   12|    1|   29|      2|      2| Atlantic|
| Mel|   31|   34|   44|   27|      3|      3| Food Town|
| Joe|   13|   41|   12|   10|      4|      4|   Jake's|
| Jill|  23|    6|    3|   45|      3|      3| Food Town|
| Joe|   33|   16|   28|   28|      2|      2| Atlantic|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```