

# **Illinois Institute of Technology**



## **CSP 554 - Big Data Technologies**

### **Project Report**

**Topic 1:** Explore use of cloud NoSQL databases in depth

### **In depth Comparison of MongoDB, DynamoDB, and Cassandra NO-SQL Databases.**

Prof. Joseph Rosen

#### **Team Members:**

Utkaarsh Vinod Bhaskarwar - A20511982  
Ashmita Gupta - A20512498  
Sujith Chandra Dara - A20528384  
Abhishek Wagh - A20518938

## **Project Outline**

This project focuses on optimizing heart disease data management through Create, Read, Update, and Delete (CRUD) operations in three prominent NoSQL databases—MongoDB, DynamoDB, and Cassandra. This project centers on an in-depth exploration and analysis of three prominent NoSQL databases aiming to comprehend their distinct use cases and application scenarios. Each database exhibits unique features catering to specific industries, and our exploration will emphasize understanding their strengths and limitations. The project will further delve into practical aspects by focusing on Create, Read, Update, and Delete (CRUD) operations on a dataset, providing hands-on experience with each database. An essential aspect of this investigation involves a comparative analysis of the performance of CRUD operations across MongoDB, DynamoDB and Cassandra shedding light on their respective efficiencies and suitability for diverse application scenarios.

## **Literature Review**

### **1. MongoDB**

MongoDB is a NoSQL database that stores data in flexible, JSON-like documents, rather than the rigid rows and columns of traditional relational databases. It is a popular choice for applications that require high scalability, performance, and flexibility[2]. MongoDB's document-oriented storage allows for easy storage and retrieval of complex data. Unlike relational databases, it does not require a predefined schema, meaning the structure of documents can evolve over time without requiring schema changes. This makes MongoDB well-suited for applications with dynamic data requirements. It also supports powerful queries using its proprietary query language, MongoDB Query Language (MQL). MQL allows for complex filtering, sorting, and aggregation of data, making it a versatile tool for data analysis and manipulation[7].

Its horizontal scalability makes it possible to handle increasing data volumes and workloads effectively. This is achieved by adding additional nodes to a cluster, allowing the database to distribute data and processing across multiple machines. Its high performance is particularly evident for read and write operations, making it well-suited for applications that require real-time data access. This performance is attributed to its efficient data storage and query execution mechanisms. MongoDB's flexibility and scalability make it a popular choice for a variety of applications, including web applications, mobile applications, real-time analytics applications, content management systems, and e-commerce applications[11]. Its ability to handle large volumes of data efficiently, its powerful query language, and its high performance make it a valuable tool for modern applications.

MongoDB redefines data storage with its NoSQL approach, utilizing flexible JSON-like documents instead of rigid relational structures[7]. This database is a top choice for applications seeking scalability, performance, and adaptability. MongoDB's document-oriented storage simplifies complex data handling, and its dynamic schema allows for evolving structures without schema changes. With its powerful query language (MQL), MongoDB facilitates intricate data analysis, making it a versatile tool for modern applications, including web and mobile apps, analytics, content management, and e-commerce.

## 2. DynamoDB

Amazon DynamoDB is a popular and fully managed NoSQL database that uses a distributed hash table (DHT) data structure. This data structure allows for efficient and consistent read and write operations across multiple nodes, ensuring high availability and low-latency access to data<sup>[4]</sup>. DynamoDB's DHT-based architecture enables automatic partitioning and distribution of data, eliminating the need for manual sharding and providing a hassle-free experience for developers. DynamoDB supports key-value and document data models, providing flexibility for developers. While DynamoDB offers advantages in scalability and flexibility, it comes with certain drawbacks<sup>[9]</sup>. The cost associated with provisioned throughput and storage, as well as the eventual consistency model.

In real-world applications, DynamoDB shines in e-commerce. Its scalability proves invaluable for handling varying traffic levels during peak seasons, ensuring a responsive and reliable shopping experience. The key-value nature and automatic partitioning of DynamoDB simplify product information retrieval<sup>[13]</sup>, exemplifying its practicality in addressing the dynamic needs of high-traffic platforms.

Amazon DynamoDB, a fully managed NoSQL database, employs a distributed hash table (DHT) structure for efficient and consistent operations across nodes<sup>[4]</sup>. Despite considerations like provisioned throughput costs and eventual consistency, DynamoDB shines in real-world applications, especially in e-commerce. Its scalability proves invaluable during peak traffic, ensuring a responsive shopping experience. DynamoDB's key-value model and automatic partitioning simplify data retrieval, exemplifying its practicality for high-traffic platforms. In essence, DynamoDB offers a powerful, hassle-free solution for developers managing diverse datasets in dynamic environments.

## 3. Cassandra

A popular NoSQL database system with strong fault tolerance and scalability is called Cassandra. The Apache Software Foundation created it, and managing massive volumes of dispersed data over numerous commodity computers is one of its main uses. Because there isn't a single point of failure, Cassandra's decentralized peer-to-peer architecture offers high availability and fault tolerance<sup>[3]</sup>. Because of its column-family-oriented data model, which enables dynamic and adaptable schema designs, it can handle a wide range of changing and diversified data structures. Because Cassandra can read and write large volumes of data with little latency, it is a popular choice for applications that need high performance and seamless scaling in highly available, distributed systems.

Cassandra's commitment to distributed consistency and durability sets it apart in the realm of NoSQL databases. Leveraging a tunable consistency model, Cassandra allows users to configure the level of consistency for their operations, striking a balance between performance and reliability<sup>[3]</sup>. Additionally, the system ensures durability by replicating data across multiple nodes, safeguarding against potential hardware failures or data center outages. This distributed approach not only enhances fault tolerance but also contributes to data resilience and reliability. Cassandra's focus on maintaining consistency and durability makes it an excellent choice for applications where data integrity is paramount, such as in financial systems, healthcare databases, and other mission-critical environments where accuracy and reliability are non-negotiable.

## Key Differences:

Features	MongoDB	DynamoDB	Cassandra
<b>Data Model</b>	<ul style="list-style-type: none"> <li>- Document oriented data model using Binary JSON(BSON)<a href="#">[7]</a></li> <li>- Allows flexible fields and supports dynamic schemas</li> </ul>	<ul style="list-style-type: none"> <li>- Key-value and document data model, using JSON-like documents.</li> <li>- Each item has a primary key that is made up of several properties.</li> </ul>	<ul style="list-style-type: none"> <li>- Wide-column store with tables and rows identified by a primary key.<a href="#">[3]</a></li> <li>- Offers a flexible schema with support for sparse columns.</li> </ul>
<b>Query Language</b>	Rich query language with rich query support, indexing, and aggregation.	Basic querying capabilities, focused on key-based operations.	CQL (Cassandra Query Language) is used, which is similar to standard SQL.
<b>Cost</b>	<ul style="list-style-type: none"> <li>- A free open-source version is available.</li> <li>- Commercial versions and managed services might come with additional charges.</li> </ul>	<ul style="list-style-type: none"> <li>- A pay-as-you-go pricing mechanism based on provisioned throughput and storage.</li> <li>- A free tier is offered for limited use.</li> </ul>	<ul style="list-style-type: none"> <li>- There is a free open-source version available.</li> <li>- Costs associated with self-hosted deployments and managed services.</li> </ul>
<b>Operating system</b>	Compatibility with a variety of operating systems such as Windows, Linux, and macOS <a href="#">[6]</a>	Serverless computing; the underlying infrastructure is abstracted from the user.	Compatibility with a variety of operating systems such as Windows, Linux, and macOS
<b>Performance</b>	Excellent read performance on huge datasets.	Consistent and predictable performance. <a href="#">[7]</a>	High write and read throughput, optimized for workloads that require a lot of writing.
<b>Primary Key</b>	The primary key is defined by the user and is based on the BSON document structure.	The user-defined primary key can be a single attribute or a composite.	The primary key is defined by the user and might be simple or composite.
<b>ACID properties</b>	Individual document ACID transactions are supported.	Ensures ACID transactions for individual items within a partition. <a href="#">[8]</a>	Limited ACID support, defaults to eventual consistency.
<b>Use Cases</b>	Versatile, suitable for various applications. <a href="#">[6]</a>	Well-suited for predictable and variable workloads.	Commonly used in scenarios requiring high write and read.
<b>Deployment</b>	It is possible to self-host or use third-party services to manage it.	AWS manages everything, there is no self-hosting alternative.	It is possible to self-host or use third-party services to manage

<b>Security</b>	<ul style="list-style-type: none"> <li>- RBAC (role-based access control).</li> <li>- Encryption at rest and in transit.</li> </ul>	<ul style="list-style-type: none"> <li>- Integration with AWS Identity and Access Management (IAM).</li> <li>- Encryption at rest and in transit.</li> </ul>	<ul style="list-style-type: none"> <li>- RBAC (role-based access control).<a href="#">[3]</a></li> <li>- Transparent data encryption.</li> </ul>
<b>Architecture</b>	<ul style="list-style-type: none"> <li>- Replica sets for high availability.</li> <li>- Master-slave architecture in sharded clusters.</li> </ul>	<ul style="list-style-type: none"> <li>- Fully managed service with a distributed multi-AZ architecture.</li> <li>- Consistent hashing for partitioning data.</li> </ul>	<ul style="list-style-type: none"> <li>- Masterless architecture with peer-to-peer node communication.</li> <li>- Data distribution across nodes using consistent hashing.</li> </ul>
<b>Indexing</b>	Allows for the creation of various index types, like compound, unique, text, and geospatial indexes. <a href="#">[7]</a>	Automatic indexing on primary key properties, including secondary index support.	Indexing is supported, including secondary indexes and custom index types. <a href="#">[3]</a>
<b>Memory features</b>	RAM is effectively used for caching frequently requested data.	AWS manages the service, details are abstracted. <a href="#">[8]</a>	Memory is used efficiently for caching frequently requested data.
<b>Consistency Model</b>	Eventual consistency is the default, although strong consistency is supported for specified operations.	Offers tunable consistency, allowing a choice between eventual consistency and strong consistency.	Defaults to eventual consistency with configurable consistency levels.

## Project Plan:

Phase	Task Description
Project Scope and Objectives	A. Define the project's scope, outlining the boundaries and goals. B. Defining the objectives of implementing CRUD operations for MongoDB, DynamoDB, and Cassandra.
Requirements Analysis	A. Identifying specific CRUD operations required for MongoDB, DynamoDB, and Cassandra. B. Defining the data model for each database, considering the unique features and requirements of each.
Technology Selection	Checking the availability of necessary drivers, SDKs, and libraries for setup of MongoDB, DynamoDB, and Cassandra
Environment Setup	A. Set up development, testing, and production environments for MongoDB, DynamoDB, and Cassandra. B. Installing and configuring the required software and dependencies for each NoSQL database.
Database Design	A. Designing the schema for MongoDB collections, DynamoDB tables, and Cassandra tables.
CRUD Operations Implementation	A. Create scripts or functions for each CRUD operation tailored to MongoDB, DynamoDB, and Cassandra. [10] B. Optimize operations to leverage the strengths of each database, considering factors such as scalability and performance.[9]
Documentation	A. Documenting the implemented CRUD operations, including usage guidelines and best practices. B. Provide comprehensive documentation for the database schemas and any configuration specifics in our final report

## DynamoDB Implementation Details

### Uploading the data on the DynamoDB

#### Step 1:

- A. Getting AWS access key and secret access key and then exporting it to form a connection with the DynamoDB

```
abhim@Abhishek MINGW64 /e/Temp
$ export AWS_ACCESS_KEY_ID=AKIA3WFYOBUNVF6EZP56
abhim@Abhishek MINGW64 /e/Temp
$ export AWS_SECRET_ACCESS_KEY=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
```

- B. Installing boto3, pandas and openpyxl

```
abhim@Abhishek MINGW64 /e/Temp
$ pip install boto3
```

```
abhim@Abhishek MINGW64 /e/Temp
$ pip install pandas
```

```
abhim@Abhishek MINGW64 /e/Temp
$ pip install openpyxl
```

- C. Executing a python code to upload a data from local system to the DynamoDB database:  
Here we are using heart\_2022\_no\_nans data from kaggle

```
abhim@Abhishek MINGW64 /e/Temp
$ python upload_data_to_dynamodb.py
Data uploaded to heart_2022_no_nans successfully!
```

- D. After successfully executing the Python code, we will be able to see the "heart\_2022\_no\_nans" data in the DynamoDB tables that we created earlier.

ws.amazon.com/dynamodbv2/home?region=us-east-1#item-explorer?operation=SCAN&table=heart\_2022\_no\_nans

[Alt+S]

DynamoDB > Explore items > heart\_2022\_no\_nans

### Tables (1)

Any tag key  
Any tag value  
Find tables by table name

heart\_2022\_no\_nans

### heart\_2022\_no\_nans

Autopreview View table details

Scan or query items

This table has more items to retrieve. To retrieve the next page of items, choose Retrieve next page.

Items returned (50)

ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	D
251	Age 55 to 59	Yes	27.46	Yes	No	N
187	Age 70 to 74	No	30.67	Yes	No	N
154	Age 70 to 74	No	23.69	Yes	Yes	Y
7	Age 80 or older	No	33.3	Yes	No	N
115	Age 65 to 69	No	34.87	Yes	No	N
286	Age 80 or older	No	19.22	Yes	No	N
361	Age 70 to 74	Yes	34.46	Yes	No	N
380	Age 80 or older	No	25.51	No	Yes	N
117	Age 65 to 69	No	26.58	Yes	No	Y

amazon.com/dynamodbv2/home?region=us-east-1#item-explorer?operation=SCAN&table=heart\_2022\_no\_nans

[Alt+S]

Find tables by table name

heart\_2022\_no\_nans

This table has more items to retrieve. To retrieve the next page of items, choose Retrieve next page.

Items returned (100)

ID (Number)	Hours	SmokerStatus	State	TetanusLast10Tdap	WeightInKilograms
287		Never smoked	Alabama	Yes, received Tdap	62.6
468		Never smoked	Alabama	No, did not receive an...	72.57
241		Former smoker	Alabama	Yes, received Tdap	65.32
339		Former smoker	Alabama	No, did not receive an...	74.84
329		Never smoked	Ala... <input type="button" value="Edit"/> <input type="button" value="Delete"/>	Yes, received Tdap	92.99
345		Former smoker	Alabama	Yes, received tetanus ...	99.79
296		Never smoked	Alabama	No, did not receive an...	62.14
75		Never smoked	Alabama	Yes, received tetanus ...	58.97
254		Never smoked	Alabama	No, did not receive an...	88.9
163		Never smoked	Alabama	No, did not receive an...	70.31
89		Never smoked	Alabama	No, did not receive an...	57.61
32		Former smoker	Alabama	Yes, received tetanus ...	81.65
131		Never smoked	Alabama	No, did not receive an...	92.99
435		Never smoked	Alabama	No, did not receive an...	67.13

Now, we will perform CRUD operations—insert, read, update, and delete using DynamoDB. Additionally, we will take note of the time it takes to execute the queries and retrieve the results.

## Insert Operations:

### Insert single row:

Here we are inserting single row with ID= 502

The screenshot shows the AWS DynamoDB console interface. On the left, there's a navigation sidebar with options like Dashboard, Tables, PartiQL editor, Backups, Exports to S3, Imports from S3, Integrations, Reserved capacity, and Settings. The main area is titled "Items returned (400)" and displays a table with 400 rows. The columns are: ID (Number), AgeCategory, AlcoholDrinkers, BMI, ChestScan, CovidPos, and D. The first few rows are as follows:

ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	D
484	Age 80 or older	No	58.23	Yes	No	N
485	Age 25 to 29	Yes	32.55	Yes	Yes	Y
486	Age 18 to 24	Yes	17.5	No	Yes	N
487	Age 70 to 74	No	23.5	No	No	N
488	Age 70 to 74	Yes	24.39	Yes	No	N
489	Age 45 to 49	Yes	29.84	Yes	No	N
490	Age 30 to 34	Yes	22.14	Yes	No	N
491	Age 50 to 54	No	36.28	No	Yes	N
492	Age 25 to 29	Yes	26.63	Yes	Yes	N
493	Age 60 to 64	No	18.13	Yes	No	N
494	Age 65 to 69	Yes	33.58	Yes	Yes	N
496	Age 45 to 49	Yes	24.14	No	Yes	N
499	Age 40 to 44	Yes	27.44	No	No	N
500	Age 55 to 59	Yes	25.82	Yes	No	N
501	Age 55 to 59	Yes	22.82	Yes	Yes	N
502	Age 55 to 59	Yes	22.82	Yes	Yes	N

### Insert multiple rows:

Here we inserted multiple rows with ID's 503,504,505

The screenshot shows a terminal window with the following output:

```
Time taken: 0.3991506099700928 seconds

abhim@Abhishek MINGW64 /e/Temp
$ python insertmany.py
3 items inserted into heart_2022_no_nans successfully!
Time taken: 0.44214558601379395 seconds
```

The screenshot shows the AWS DynamoDB console interface. On the left, there's a sidebar with navigation links like Dashboard, Tables, Explore items, and DAX. The main area is titled 'Scan or query items' and shows a table with 500 items. The columns include ID (Number), AgeCategory, AlcoholDrinkers, BMI, ChestScan, CovidPos, and D. The first few rows of data are:

ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	D
505	Age 55 to 59	Yes	22.82	Yes	Yes	N
504	Age 55 to 59	Yes	22.82	Yes	Yes	N
503	Age 55 to 59	Yes	22.82	Yes	Yes	N
502	Age 55 to 59	Yes	22.82	Yes	Yes	N
501	Age 55 to 59	Yes	22.82	Yes	Yes	N
500	Age 55 to 59	Yes	25.82	Yes	No	N
499	Age 40 to 44	Yes	27.44	No	No	N
498	Age 70 to 74	No	29.01	Yes	No	N
497	Age 30 to 34	No	40.39	No	Yes	N
496	Age 45 to 49	Yes	24.14	No	Yes	N
495	Age 75 to 79	No	21.13	Yes	No	Y

## Read Operations:

### Read single row:

Here we are reading a single row with ID = 501.

```
abhim@Abhishek MINGW64 /e/Temp
$ python read1.py
Item read from heart_2022_no_nans: {'HadDiabetes': 'Yes', 'HeightInMeters': '1.68', 'CovidPos': 'Yes', 'LastCheckupTime': 'Within past year (anytime less than 12 months ago)', 'SleepHours': '9', 'ChestScan': 'Yes', 'HighRiskLastYear': 'No', 'HadDepressiveDisorder': 'No', 'AlcoholDrinkers': 'Yes', 'HadHeartAttack': 'No', 'WeightInKilograms': '77.25', 'HadAsthma': 'Yes', 'DeafOrHardOfHearing': 'No', 'ID': Decimal('501'), 'EcigaretteUsage': 'Never Used', 'RaceEthnicityCategory': 'White only, Non-Hispanic', 'HadAngina': 'No', 'DifficultyConcentrating': 'Yes', 'HadArthritis': 'No', 'HadCOPD': 'Yes', 'Sex': 'Female', 'PhysicalActivities': 'Yes', 'HadStroke': 'Yes', 'DifficultyWalking': 'No', 'SmokerStatus': 'Former smoker', 'TetanusLast10Tdap': 'Yes, received Tdap', 'HadSkinCancer': 'No', 'MentalHealthDays': '0', 'FluVaxLast12': 'No', 'GeneralHealth': 'Very Good', 'State': 'Illinois', 'AgeCategory': 'Age 55 to 59', 'PhysicalHealthDays': '4', 'PneumoVaxEver': 'Yes', 'HadKidneyDisease': 'Yes', 'HIVTesting': 'No', 'BMI': '22.82'}
Time taken for read operation: 0.4097745418548584 seconds

abhim@Abhishek MINGW64 /e/Temp
$
```

## Reading all the rows in from the DynamoDB database:

```
abhim@Abhishek MINGW64 /e/Temp
$ python readall.py
Time taken for scan operation: 0.6462013721466064 seconds
```

## **Update operations:**

### Update specific row:

Here we are updating the GeneralHealth of ID =501

```
MINGW64:/e/Temp
abhim@Abhishek MINGW64 /e/Temp
$ python update1.py
item with key {'ID': 501} updated successfully!
Time taken for update operation: 0.40413951873779297 seconds
abhim@Abhishek MINGW64 /e/Temp
$ |
```

The screenshot shows the AWS DynamoDB console interface. On the left, there's a navigation sidebar with options like Dashboard, Tables, Update settings, Explore items, PartiQL editor, Backups, Exports to S3, Imports from S3, Integrations, Reserved capacity, and Settings. Under the 'Tables' section, it shows 'Tables (1)'. The main area is titled 'heart\_2022\_no\_nans'. It displays a table with the following data:

ID (Number)	FluVaxLast12	GeneralHealth	HadAngina	HadArthritis	HadAsth...
501	No	Bad	No	No	Yes
401	Yes	Excellent	No	Yes	No
273	Yes	Excellent	No	No	No
145	Yes	Excellent	No	No	No
467	No	Excellent	No	No	No
187	No	Fair	No	Yes	No
154	Yes	Fair	No	No	No
117	No	Fair	No	No	No
47	Yes	Fair	No	Yes	No

### Update multiple rows:

Here we are updating the GeneralHealth of IDs 501,502,503

```
abhim@Abhishek MINGW64 /e/Temp
$ python updatemany.py
Items updated successfully!
Time taken for update operations: 0.5311610698699951 seconds
```

The screenshot shows the AWS DynamoDB console. On the left, the navigation pane includes 'Dashboard', 'Tables', 'Explore items' (which is selected), 'PartiQL editor', 'Backups', 'Exports to S3', 'Imports from S3', 'Integrations', 'Reserved capacity', and 'Settings'. Under 'DAX', there are 'Clusters', 'Subnet groups', 'Parameter groups', and 'Events'. The main area displays the 'heart\_2022\_no\_nans' table, which has 1 item. A message at the top right says 'Completed. Read capacity units consumed: 0.5'. Below it is a table titled 'Items returned (505)' with columns: ID (Number), FluVaxLast12, GeneralHealth, HadAngina, HadArthritis, HadAsthma, and others. The first five rows are shown:

ID (Number)	FluVaxLast12	GeneralHealth	HadAngina	HadArthritis	HadAsthma
501	No	Good	No	No	Yes
502	No	Not Bad	No	No	Yes
503	No	Really Bad	No	No	Yes
504	No	Very Good	No	No	Yes

## Delete Operations:

### Delete single row:

Here we are deleting a row with ID 501.

```
abhim@Abhishek MINGW64 /e(Temp
$ python deletespecific.py
Item with key {'ID': 501} deleted successfully!
Time taken for delete operation: 0.3968343734741211 seconds
```

Here, we can see the result in DynamoDB that the row with the ID 501 has been successfully deleted.

The screenshot shows the AWS DynamoDB console. The table 'heart\_2022\_no\_nans' now has 504 items. The table view shows the following data:

ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	Death
502	Age 55 to 59	Yes	22.82	Yes	Yes	No
503	Age 55 to 59	Yes	22.82	Yes	Yes	No
504	Age 55 to 59	Yes	22.82	Yes	Yes	No
505	Age 55 to 59	Yes	22.82	Yes	Yes	No

## Delete Multiple rows:

Here, we are deleting multiple rows from DynamoDB with the IDs 1, 502, 503, and 504.

```
abhim@Abhishek MINGW64 /e/Temp
$ python deletemany.py
Items with keys [{'ID': 1}, {'ID': 502}, {'ID': 503}, {'ID': 504}] deleted successfully!
Time taken for delete operations: 0.5258674621582031 seconds
```

Here in the DynamoDB table, we can see that the row with ID 1 has been deleted.

Items returned (500)								
	ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	D	
<input type="checkbox"/>	<a href="#">17</a>	Age 65 to 69	No	25.54	Yes	No	N	
<input type="checkbox"/>	<a href="#">16</a>	Age 80 or older	No	28.28	Yes	Yes	N	
<input type="checkbox"/>	<a href="#">15</a>	Age 70 to 74	No	27.26	Yes	Yes	Y	
<input type="checkbox"/>	<a href="#">14</a>	Age 60 to 64	No	35.15	No	No	N	
<input type="checkbox"/>	<a href="#">13</a>	Age 60 to 64	No	32.74	Yes	No	N	
<input type="checkbox"/>	<a href="#">12</a>	Age 60 to 64	No	46.87	No	Yes	N	
<input type="checkbox"/>	<a href="#">11</a>	Age 80 or older	Yes	36.62	Yes	No	N	
<input type="checkbox"/>	<a href="#">10</a>	Age 75 to 79	No	22.6	Yes	No	N	
<input type="checkbox"/>	<a href="#">9</a>	Age 40 to 44	No	26.94	Yes	Yes	Y	
<input type="checkbox"/>	<a href="#">8</a>	Age 75 to 79	No	24.37	Yes	Yes	N	
<input type="checkbox"/>	<a href="#">7</a>	Age 80 or older	No	33.3	Yes	No	N	
<input type="checkbox"/>	<a href="#">6</a>	Age 50 to 54	Yes	34.96	Yes	No	N	
<input type="checkbox"/>	<a href="#">5</a>	Age 80 or older	No	33.07	No	No	N	
<input type="checkbox"/>	<a href="#">4</a>	Age 80 or older	No	31.32	No	Yes	N	
<input type="checkbox"/>	<a href="#">3</a>	Age 75 to 79	Yes	31.66	Yes	Yes	N	
<input type="checkbox"/>	<a href="#">2</a>	Age 70 to 74	No	30.13	No	No	N	

And, here in the DynamoDB table, we can see that the rows with IDs 502, 503 and 504 have been deleted.

Items returned (500)

Actions ▾ Create item

ID (Number)	AgeCategory	AlcoholDrinkers	BMI	ChestScan	CovidPos	D
<a href="#">505</a>	Age 55 to 59	Yes	22.82	Yes	Yes	N
<a href="#">500</a>	Age 55 to 59	Yes	25.82	Yes	No	N
<a href="#">499</a>	Age 40 to 44	Yes	27.44	No	No	N
<a href="#">498</a>	Age 70 to 74	No	29.01	Yes	No	N
<a href="#">497</a>	Age 30 to 34	No	40.39	No	Yes	N
<a href="#">496</a>	Age 45 to 49	Yes	24.14	No	Yes	N
<a href="#">495</a>	Age 75 to 79	No	21.13	Yes	No	Y
<a href="#">494</a>	Age 65 to 69	Yes	33.58	Yes	Yes	N

## Other Operations:

Here, we are performing the min, max, and average operations to calculate BMI from the DynamoDB table.

### Minimum value:

Here we are deriving the minimum value of BMI

```
abhim@Abhishek MINGW64 /e/Temp
$ python min.py
Minimum BMI: 15.66
Time taken: 0.5072319507598877 seconds
```

### Maximum value:

Here we are deriving the maximum value of BMI

```
abhim@Abhishek MINGW64 /e/Temp
$ python max.py
Maximum BMI: 91.55
Time taken: 0.4986891746520996 seconds
```

### Average value:

Here we are deriving the average value of BMI

```
abhim@Abhishek MINGW64 /e/Temp
$ python avg.py
Average BMI: 29.15656
Time taken: 0.4435694217681885 seconds
```

## Latency of CRUD operations:

Finally, we perform all CRUD operations in a single Python code to calculate the latency of each operation.

```
abhim@Abhishek MINGW64 /e/Temp
$ python latency.py
Create Result: {'ResponseMetadata': {'RequestId': '11VKVP1MTRVNSONHIBVG5MHPFJVV4KQNS05AEMVJF66Q9ASUAAJG', 'HTTPStatusCode': 200, 'HTTPHeaders': {'server': 'Server', 'date': 'Sun, 03 Dec 2023 18:23:48 GMT', 'content-type': 'application/x-amz-json-1.0', 'content-length': '2', 'connection': 'keep-alive', 'x-amzn-requestid': '11VKVP1MTRVNSONHIBVG5MHPFJVV4KQNS05AEMVJF66Q9ASUAAJG', 'x-amz-crc32': '2745614147'}, 'RetryAttempts': 0}}
Create Latency: 0.4214816093444824 seconds
Read Result: {'Item': {'HadDiabetes': 'Yes', 'HeightInMeters': '1.68', 'CovidPos': 'Yes', 'LastCheckupTime': 'Within past year (anytime less than 12 months ago)', 'SleepHours': '9', 'ChestScan': 'Yes', 'HighRiskLastYear': 'No', 'HadDepressiveDisorder': 'No', 'AlcoholDrinkers': 'Yes', 'HadHeartAttack': 'No', 'WeightInKilograms': '77.25', 'HadAsthma': 'Yes', 'DeafOrHardOfHearing': 'No', 'ID': Decimal('506'), 'ECigaretteUsage': 'Never Used', 'RaceEthnicityCategory': 'White only, Non-Hispanic', 'HadAngina': 'No', 'DifficultyConcentrating': 'Yes', 'HadArthritis': 'No', 'HadCOPD': 'Yes', 'Sex': 'Male', 'PhysicalActivities': 'Yes', 'HadStroke': 'Yes', 'DifficultyWalking': 'No', 'SmokerStatus': 'Former smoker', 'TetanusLast10Tdap': 'Yes, received Tdap', 'HadSkinCancer': 'No', 'MentalHealthDays': '0', 'FluVaxLast12': 'No', 'GeneralHealth': 'Very Good', 'State': 'Illinois', 'AgeCategory': 'Age 55 to 59', 'PhysicalHealthDays': '4', 'PneumoVaxEver': 'Yes', 'HadKidneyDisease': 'Yes', 'HIVTesting': 'No', 'BMI': '22.82'}, 'ResponseMetadata': {'RequestId': '8PP8KM51DIRQDDAO5FJ5CFPENV4KQNS05AEMVJF66Q9ASUAAJG', 'HTTPStatusCode': 200, 'HTTPHeaders': {'server': 'Server', 'date': 'Sun, 03 Dec 2023 18:23:48 GMT', 'content-type': 'application/x-amz-json-1.0', 'content-length': '1148', 'connection': 'keep-alive', 'x-amzn-requestid': '8PP8KM51DIRQDDAO5FJ5CFPENV4KQNS05AEMVJF66Q9ASUAAJG', 'x-amz-crc32': '2399451196'}, 'RetryAttempts': 0}}
Read Latency: 0.05133509635925293 seconds
Update Result: {'ResponseMetadata': {'RequestId': 'NBC2GK003SL7QF9HR69ENB7S5VV4KQNS05AEMVJF66Q9ASUAAJG', 'HTTPStatusCode': 200, 'HTTPHeaders': {'server': 'Server', 'date': 'Sun, 03 Dec 2023 18:23:48 GMT', 'content-type': 'application/x-amz-json-1.0', 'content-length': '2', 'connection': 'keep-alive', 'x-amzn-requestid': 'NBC2GK003SL7QF9HR69ENB7S5VV4KQNS05AEMVJF66Q9ASUAAJG', 'x-amz-crc32': '2745614147'}, 'RetryAttempts': 0}}
Update Latency: 0.03980898857116699 seconds
Delete Result: {'ResponseMetadata': {'RequestId': '38NL2ENUMONHAT6APCGH03JGPBV4KQNS05AEMVJF66Q9ASUAAJG', 'HTTPStatusCode': 200, 'HTTPHeaders': {'server': 'Server', 'date': 'Sun, 03 Dec 2023 18:23:48 GMT', 'content-type': 'application/x-amz-json-1.0', 'content-length': '2', 'connection': 'keep-alive', 'x-amzn-requestid': '38NL2ENUMONHAT6APCGH03JGPBV4KQNS05AEMVJF66Q9ASUAAJG', 'x-amz-crc32': '2745614147'}, 'RetryAttempts': 0}}
Delete Latency: 0.043770551681518555 seconds
```

## MongoDB Implementation Details

Now, we will perform CRUD operations—insert, read, update, and delete using MongoDB. Additionally, we will take note of the time it takes to execute the queries and retrieve the results.

### **Insert Operations:**

#### Insert single row:

Here we are inserting single row with ID= 502

```
"State" : "Ohio",
"Sex" : "Male",
"GeneralHealth" : "Very Good",
"PhysicalHealthDays" : "4",
"MentalHealthDays" : "0",
"LastCheckupTime" : "Within past year (anytime less than 12 months ago)",
"PhysicalActivities" : "Yes",
"SleepHours" : "9",
"HadHeartAttack" : "No",
"HadAngina" : "No",
"HadStroke" : "Yes",
"HadAsthma" : "Yes",
"HadSkinCancer" : "No",
"HadCOPD" : "Yes",
"HadDepressiveDisorder" : "No",
"HadKidneyDisease" : "Yes",
"HadArthritis" : "No",
"HadDiabetes" : "Yes",
"DeafOrHardOfHearing" : "No",
"DifficultyConcentrating" : "Yes",
"DifficultyWalking" : "No",
"SmokerStatus" : "Former smoker",
"ECigaretteUsage" : "Never Used",
"ChestScan" : "Yes",
"RaceEthnicityCategory" : "White only, Non-Hispanic",
"AgeCategory" : "Age 55 to 59",
"HeightInMeters" : "1.68",
"WeightInKilograms" : "77.25",
"BMI" : "22.82",
"AlcoholDrinkers" : "Yes",
"HIVTesting" : "No",
"FluVaxLast12" : "No",
"PneumoVaxEver" : "Yes",
"TetanusLast10Tdap" : "Yes, received Tdap",
"HighRiskLastYear" : "No",
"CovidPos" : "Yes",
"ID" : 502
```

```
> Load('./insertone.js');
Time taken 610.531 milliseconds
```

#### Insert multiple rows:

Here we inserted multiple rows with ID's 503,504,505

```
> db.heartcoll.insertMany = [ { 'State': 'Illinois', 'Sex': 'Female', 'GeneralHealth': 'Very Good', 'PhysicalHealthDays': '4', 'MentalHealthDays': '0', 'LastCheckupTime': 'Within past year (anytime less than 12 months ago)', 'PhysicalActivities': 'Yes', 'SleepHours': '9', 'HadHeartAttack': 'No', 'HadAngina': 'No', 'HadStroke': 'Yes', 'HadSkinCancer': 'No', 'HadCOPD': 'Yes', 'HadDepressiveDisorder': 'No', 'HadKidneyDisease': 'Yes', 'HadArthritis': 'No', 'HadDiabetes': 'Yes', 'DeafOrHardOfHearing': 'No', 'ChestScan': 'Yes', 'RaceEthnicityCategory': 'White only, Non-Hispanic', 'AgeCategory': 'Age 55 to 59', 'HeightInMeters': '1.68', 'WeightInKilograms': '77.25', 'BMI': '22.82', 'AlcoholDrinkers': 'Yes', 'HIVTesting': 'No', 'FluVaxLast12': 'No', 'PneumoVaxEver': 'Yes', 'TetanusLast10Tdap': 'Yes, received Tdap', 'HighRiskLastYear': 'No', 'CovidPos': 'Yes', 'ID': 503}, { 'State': 'Ohio', 'Sex': 'Male', 'GeneralHealth': 'Very Good', 'PhysicalHealthDays': '4', 'MentalHealthDays': '0', 'LastCheckupTime': 'Within past year (anytime less than 12 months ago)', 'PhysicalActivities': 'Yes', 'SleepHours': '9', 'HadHeartAttack': 'No', 'HadAngina': 'No', 'HadStroke': 'Yes', 'HadSkinCancer': 'Yes', 'HadCOPD': 'Yes', 'HadDepressiveDisorder': 'No', 'HadKidneyDisease': 'Yes', 'HadArthritis': 'No', 'HadDiabetes': 'Yes', 'DeafOrHardOfHearing': 'No', 'ChestScan': 'Yes', 'RaceEthnicityCategory': 'White only, Non-Hispanic', 'AgeCategory': 'Age 55 to 59', 'HeightInMeters': '1.68', 'WeightInKilograms': '77.25', 'BMI': '22.82', 'AlcoholDrinkers': 'Yes', 'HIVTesting': 'Yes', 'FluVaxLast12': 'No', 'PneumoVaxEver': 'Yes', 'TetanusLast10Tdap': 'Yes, received Tdap', 'HighRiskLastYear': 'No', 'CovidPos': 'Yes', 'ID': 503} ]
```

```
> load('./insertmany.js');
Time taken 431.682 milliseconds
```

## Read Operations:

### Read single row:

Here we are reading a single row with ID = 500.

```
> db.heartcoll.find({ID:500});
[ { "_id": ObjectId("6570db832fbadde835fe868a"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Fair", "PhysicalHealthDays": 14, "MentalHealthDays": 30, "LastCheckupTime": "Within past 2 years (1 year but less than 2 years ago)", "PhysicalActivities": "No", "SleepHours": 6, "HadHeartAttack": "No", "HadAngina": "Yes", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "Yes", "HadDepressiveDisorder": "Yes", "HadKidneyDisease": "No", "HadArthritis": "No", "HadDiabetes": "No", "DeafOrHardOfHearing": "No", "ChestScan": "Not at all (right now)", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 55 to 59", "HeightInMeters": 1.68, "WeightInKilograms": 77.25, "BMI": "22.82", "AlcoholDrinkers": "Yes", "HIVTesting": "No", "FluVaxLast12": "No", "PneumoVaxEver": "Yes", "TetanusLast10Tdap": "Yes, received Tdap", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 500 } ]
```

```
> load('./readOne.js');
Time taken 508.644 milliseconds
```

### Read all

Here we are reading all rows in from MongoDB database

```
> hadoop@ip-172-31-52-33:~ 
> use heart_disease
switched to db heart_disease
> db.heartcoll.find();
[ { "_id": ObjectId("6570db832fbadde835fe8498"), "State": "Alabama", "Sex": "Male", "GeneralHealth": "Very good", "PhysicalHealthDays": 0, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 6, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "Yes", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Former smoker", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "No", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 55 to 59", "HeightInMeters": 1.78, "WeightInKilograms": 95.25, "BMI": "30.13", "AlcoholDrinkers": "No", "HIVTesting": "No", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received tetanus shot but not sure what type", "HighRiskLastYear": "No", "CovidPos": "No", "ID": 2 }, { "_id": ObjectId("6570db832fbadde835fe8499"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Good", "PhysicalHealthDays": 3, "MentalHealthDays": 15, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 5, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "No", "HadArthritis": "Yes", "HadDiabetes": "No", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 70 to 74", "HeightInMeters": 1.78, "WeightInKilograms": 95.25, "BMI": "30.13", "AlcoholDrinkers": "No", "HIVTesting": "No", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received tetanus shot but not sure what type", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 2 }, { "_id": ObjectId("6570db832fbadde835fe849a"), "State": "Alabama", "Sex": "Male", "GeneralHealth": "Good", "PhysicalHealthDays": 3, "MentalHealthDays": 15, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 5, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "No", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 70 to 74", "HeightInMeters": 1.78, "WeightInKilograms": 95.25, "BMI": "30.13", "AlcoholDrinkers": "No", "HIVTesting": "No", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received tetanus shot but not sure what type", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 2 }, { "_id": ObjectId("6570db832fbadde835fe849b"), "State": "Alabama", "Sex": "Male", "GeneralHealth": "Good", "PhysicalHealthDays": 0, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 7, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "No", "HadArthritis": "No", "HadDiabetes": "No", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 75 to 79", "HeightInMeters": 1.78, "WeightInKilograms": 95.25, "BMI": "30.13", "AlcoholDrinkers": "No", "HIVTesting": "No", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received tetanus shot but not sure what type", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 2 }, { "_id": ObjectId("6570db832fbadde835fe849c"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Fair", "PhysicalHealthDays": 5, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 9, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "Yes", "HadCOPD": "No", "HadDepressiveDisorder": "Yes", "HadKidneyDisease": "Yes", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "Yes", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 80 or older", "HeightInMeters": 1.7, "WeightInKilograms": 95.25, "BMI": "33.07", "AlcoholDrinkers": "No", "HIVTesting": "No", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received tetanus shot but not sure what type", "HighRiskLastYear": "Yes", "CovidPos": "Yes", "ID": 2 }, { "_id": ObjectId("6570db832fbadde835fe849d"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Very good", "PhysicalHealthDays": 4, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 9, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "No", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Former smoker", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 65 to 69", "HeightInMeters": 1.6, "WeightInKilograms": 71.67, "BMI": "34.96", "AlcoholDrinkers": "Yes", "HIVTesting": "Yes", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received Tdap", "HighRiskLastYear": "No", "CovidPos": "No", "ID": 1 }, { "_id": ObjectId("6570db832fbadde835fe849e"), "State": "Alabama", "Sex": "Male", "GeneralHealth": "Good", "PhysicalHealthDays": 0, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 9, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "No", "HadCOPD": "No", "HadDepressiveDisorder": "No", "HadKidneyDisease": "No", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "No", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 65 to 69", "HeightInMeters": 1.6, "WeightInKilograms": 71.67, "BMI": "34.96", "AlcoholDrinkers": "Yes", "HIVTesting": "Yes", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received Tdap", "HighRiskLastYear": "No", "CovidPos": "No", "ID": 1 }, { "_id": ObjectId("6570db832fbadde835fe849f"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Fair", "PhysicalHealthDays": 5, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 9, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "Yes", "HadCOPD": "No", "HadDepressiveDisorder": "Yes", "HadKidneyDisease": "Yes", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "Yes", "SmokerStatus": "Never smoked", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 75 to 79", "HeightInMeters": 1.7, "WeightInKilograms": 90.72, "BMI": "31.99", "AlcoholDrinkers": "No", "HIVTesting": "Yes", "FluVaxLast12": "Yes", "PneumoVaxEver": "Yes, received Tdap", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 1 }, { "_id": ObjectId("6570db832fbadde835fe849g"), "State": "Alabama", "Sex": "Female", "GeneralHealth": "Very good", "PhysicalHealthDays": 0, "MentalHealthDays": 0, "LastCheckupTime": "Within past year (anytime less than 12 months ago)", "PhysicalActivities": "Yes", "SleepHours": 8, "HadHeartAttack": "No", "HadAngina": "No", "HadStroke": "No", "HadAsthma": "No", "HadSkinCancer": "Yes", "HadCOPD": "No", "HadDepressiveDisorder": "Yes", "HadKidneyDisease": "Yes", "HadArthritis": "Yes", "HadDiabetes": "Yes", "DeafOrHardOfHearing": "No", "DifficultyConcentrating": "No", "DifficultyWalking": "Yes", "SmokerStatus": "Former smoker", "ECigaretteUsage": "Never used e-cigarettes in my entire life", "ChestScan": "Yes", "RaceEthnicityCategory": "White only, Non-Hispanic", "AgeCategory": "Age 75 to 79", "HeightInMeters": 1.85, "WeightInKilograms": 108.8, "BMI": "31.66", "AlcoholDrinkers": "Yes", "HIVTesting": "Yes", "FluVaxLast12": "No", "PneumoVaxEver": "Yes, received Tdap", "HighRiskLastYear": "No", "CovidPos": "Yes", "ID": 1 } ]
```

```
> load('./readall.js');
Time taken 428.372 milliseconds
true
>
```

## Update operations:

### Update specific row:

Here we are updating the ID attribute where the ID = 500.

```
> db.heartcoll.update({ ID: 500 }, { $set: { ID: 510 } });
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> |
```

```
> load('./updateone.js');
Time taken 682.731 milliseconds
true
>
```

### Update multiple rows:

Here we are updating the ID of IDs 101, 209, 459

```
> db.heartcoll.bulkWrite([
...   {
...     updateOne: {
...       filter: { ID: 101 },
...       update: { $set: { ID: 102 } }
...     }
...   },
...   {
...     updateOne: {
...       filter: { ID: 202 },
...       update: { $set: { ID: 203 } }
...     }
...   },
...   {
...     updateOne: {
...       filter: { ID: 459 },
...       update: { $set: { ID: 450 } }
...     }
...   }
... ]);
{
  "acknowledged" : true,
  "deletedCount" : 0,
  "insertedCount" : 0,
  "matchedCount" : 3,
  "upsertedCount" : 0,
  "insertedIds" : [
    ...
  ],
  "upsertedIds" : [
    ...
  ]
}
```

```
> load('./updatemany.js');
Time taken 772.134 milliseconds
true
```

## DeleteOperations:

### Delete single row:

Here we are deleting a row with ID 10.

```
> db.heartcoll.remove({ ID: 10 })
WriteResult({ "nRemoved" : 1 })
>
```

```
hadoop@ip-172-31-52-33:~ 
> load('./deleteone.js');
Time taken 221.431 milliseconds
true
>
```

### Delete Multiple rows:

Here, we are deleting multiple rows from MongoDB with the IDs 1, 28, 36, 104 and 443.

```
> db.heartcoll.remove({ ID : { $in: [1, 28, 36, 104, 443] } }, { justOne: false })
WriteResult({ "nRemoved" : 5 })
>

> load('./deletemany.js');
Time taken 311.431 milliseconds
true
```

## Other Operations:

Here, we are performing the min, max, and average operations to calculate BMI from the MongoDB table.

### Minimum value:

Here we are deriving the minimum value of BMI

```
> db.heartcoll.aggregate([
...   {
...     $group: {
...       _id: null,
...       minBMI: { $min: "$BMI" }
...     }
...   }
... ]);
{ "_id" : null, "minBMI" : 15.66 }
>
```

```
> load('./min.js');
Time taken 694.761 milliseconds
true
```

### Maximum value:

Here we are deriving the maximum value of BMI

```
> db.heartcoll.aggregate([
...   {
...     $group: {
...       _id: null,
...       maxBMI: { $max: "$BMI" }
...     }
...   }
... ]);
{ "_id" : null, "maxBMI" : 91.55 }
>
```

```
> load('./max.js');
Time taken 558.989 milliseconds
true
>
```

### Average value:

Here we are deriving the average value of BMI

```
> db.heartcoll.aggregate([
...   {
...     $group: {
...       _id: null,
...       avgBMI: { $avg: "$BMI" }
...     }
...   }
... ]);
{ "_id" : null, "avgBMI" : 29.180526315789475 }
>
```

```
> load('./avg.js');
Time taken 579.237 milliseconds
true
>
```

### Comparison Table:

	DynamoDB		MongoDB	
Operations	Time taken ( In Seconds)	Time taken ( In milliseconds)	Time taken (In Seconds)	Time taken (In milliseconds)
Inserts single row	0.399 secs	399.150 ms	0.610 secs	610.531 ms
Insert multiple rows	0.442 secs	442.145 ms	0.431 secs	431.682 ms
Read single row	0.409 secs	409.774 ms	0.508 secs	508.644 ms
Read all rows	0.646 secs	646.201 ms	0.428 secs	428.372 ms
Update specific row	0.404 secs	404.139 ms	0.682 secs	682.731 ms
Update multiple rows	0.531 secs	531.161 ms	0.772 secs	772.134 ms
Delete single row	0.396 secs	396.834 ms	0.221 secs	221.431 ms
Delete multiple rows	0.525 secs	525.867 ms	0.311 sec	311.431 ms
Minimum value	0.507 secs	507.231 ms	0.694 secs	694.761 ms
Maximum value	0.498 secs	498 ms	0.558 secs	558.989 ms
Average value	0.443 secs	433.569 ms	0.579 secs	579.237 ms

## **Analysis of results:**

### **Insert Operations:**

When it comes to single-row inserts, DynamoDB outshines MongoDB, completing the operation in 0.399 seconds compared to MongoDB's 0.610 seconds. However, MongoDB shows a slight advantage in inserting multiple rows, taking 0.431 seconds compared to DynamoDB's 0.442 seconds.

### **Read Operations:**

DynamoDB demonstrates faster performance in reading a single row, clocking in at 0.409 seconds, while MongoDB takes 0.508 seconds. Surprisingly, MongoDB excels in reading all rows, surpassing DynamoDB significantly with a time of 0.428 seconds compared to DynamoDB's 0.646 seconds.

### **Update Operations:**

For update operations, DynamoDB performs better in both specific row updates (0.404 seconds) and multiple row updates (0.531 seconds) compared to MongoDB, which takes 0.682 seconds and 0.772 seconds, respectively.

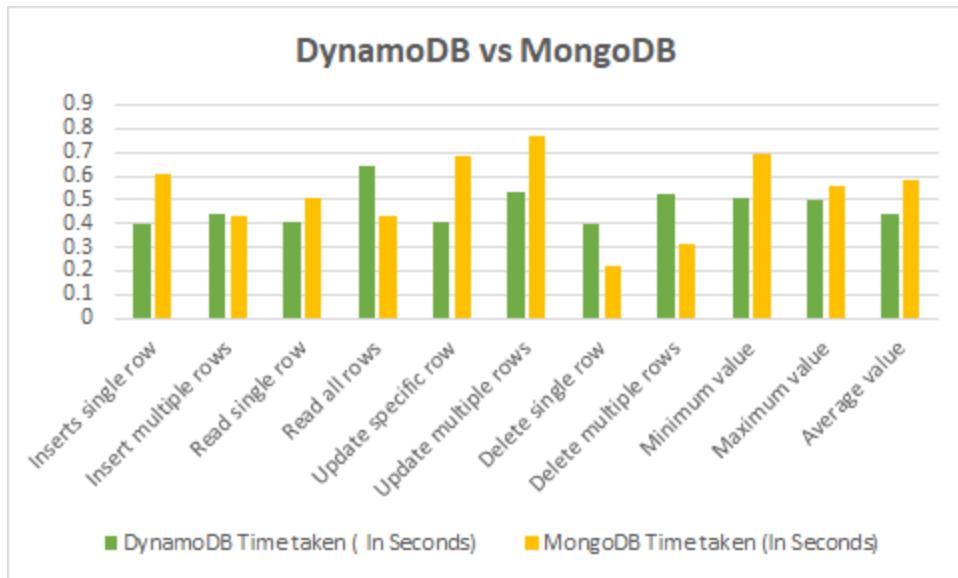
### **Delete Operations:**

MongoDB takes the lead in single-row deletions, completing the operation in 0.221 seconds, while DynamoDB requires 0.396 seconds. MongoDB also proves more efficient in deleting multiple rows, finishing the task in 0.311 seconds, compared to DynamoDB's 0.525 seconds.

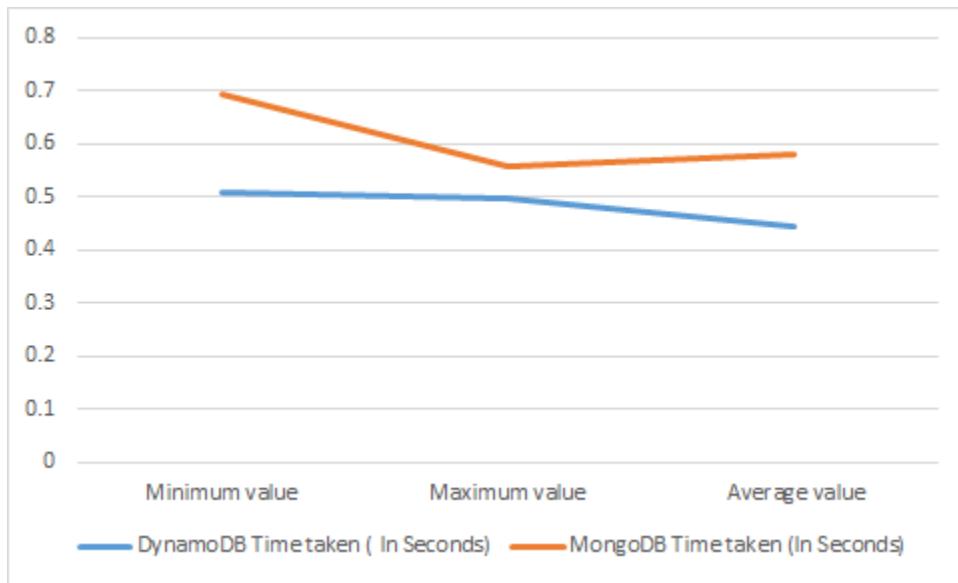
### **Aggregated Operations:**

Examining aggregated results, DynamoDB consistently outperforms MongoDB in terms of minimum time (0.507 seconds vs. 0.694 seconds), maximum time (0.498 seconds vs. 0.558 seconds), and average time (0.443 seconds vs. 0.579 seconds).

Further, we can see from the graph below that DynamoDB performs better than MongoDB for most of the operations:



Graph to show that DynamoDB performs better than MongoDB for Min, Max and Average operations:



## **Conclusion:**

After performing a set of CRUD operations on a dataset for heart disease, DynamoDB emerged as the frontrunner, showcasing superior performance across various CRUD operations. Its efficiency in tasks such as single-row inserts, updates, and deletes is notable. However, MongoDB demonstrates its own strengths, particularly in scenarios involving bulk inserts and single-row deletions, where it exhibits slightly faster response times in milliseconds.

When we examined the broader spectrum of CRUD operations, the detailed comparison underscores that both databases present nuanced trade-offs. DynamoDB excels in write-intensive tasks, while MongoDB proves adept at handling read-intensive operations. Further, we can conclude that the selection between the two should be driven by the specific demands of the use case, accounting for factors like dataset size, query complexity, and database configuration.

In conclusion, the optimal choice hinges on a careful consideration of these factors and an alignment with the application's priorities. While DynamoDB and MongoDB each have their merits, the right decision depends on the nuanced requirements of the given scenario.

## **Milestones:**

TASK	ASSIGNED TO	ETA
<b>Data setup - DynamoDB</b>	Utkarsh	12th Nov
<b>Data setup - MongoDB</b>	Ashmita	12th Nov
<b>CRUD - DynamoDB</b>	Abhishek	18th Nov
<b>CRUD - MongoDB</b>	Sujith	19th Nov
<b>Cassandra</b>	Utkarsh, Abhishek	21st Nov
<b>Performance comparison</b>	Ashmita, Sujith	25th Nov
<b>Final report</b>	Utkarsh, Abhishek, Ashmita, Sujith	30th Nov

## **References:**

1. "Dataset Link"  
<https://www.kaggle.com/datasets/kamilpytlak/personal-key-indicators-of-heart-disease>
2. "Document Database - MongoDB"  
<https://www.mongodb.com/document-databases>
3. "Document Database - Cassandra"  
<https://cassandra.apache.org/>
4. "Document Database - Amazon DynamoDB"  
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>
5. "Comparison of Databases"  
<https://ieeexplore.ieee.org/document/8284494>
6. "Comparison of Databases"  
<https://db-engines.com/en/system/Amazon+DynamoDB%3BCassandra%3BMongoDB>
7. "Comparison of MongoDB and DynamoDB"  
<https://www.mongodb.com/compare/mongodb-dynamodb>
8. "Comparison of MongoDB and Cassandra"  
<https://www.knowledgehut.com/blog/data-science/cassandra-vs-mongodb#cassandra-vs-mongodb-differences%C2%A0>
9. "Amazon DynamoDB: A Scalable, Predictably Performant, and Fully Managed NoSQL Database Service"  
<https://www.usenix.org/conference/atc22/presentation/elhemali>
10. Salimon, S. Get Started with MongoDB CRUD Operations in Python: 3 Easy Steps. Learn | Hevo.  
<https://hevodata.com/learn/mongodb-crud-operations-in-python/>
11. F. (2022, January 6). MongoDB vs. DynamoDB - Fintelics. Medium.  
<https://fintelics.medium.com/mongodb-vs-dynamodb-eff74708201e>
12. Choudhary, D, CRUD OPERATIONS FOR AWS DynamoDB USING PYTHON BOTO3 SCRIPT. Dheeraj Choudhary's Blog.  
<https://dheeraj3choudhary.com/crud-operations-for-aws-dynamodb-using-python-boto3-script>
13. TravisMedia, DynamoDB Crud Examples With Boto3 and Python.  
<https://travis.media/dynamodb-crud-examples-with-boto3-python/>
14. Nick, P. (2019, October 31). MongoDB vs. DocumentDB: Which Is Right for You? GetStream.io. Retrieved from <https://getstream.io/blog/mongodb-vs-documentdb-which-is-right-for-you/>
15. Pandas Development Team. (n.d.). pandas documentation. Retrieved from  
<https://pandas.pydata.org/pandas-docs/stable/index.html>